

HashedIn by Deloitte Data Engineer Interview Guide – Experienced 3+

Technical Round 1

1. Add a new column with manager names for each employee

Given a table structure with Id, Name, and Manager_Id, use a **self-join** to find the manager's name.

SQL Query:

```
SELECT e.Id, e.Name, e.Manager_Id, m.Name AS Manager_Name
FROM Employee e
LEFT JOIN Employee m ON e.Manager_Id = m.Id;
```

2. Identify who is a manager and who is not

To find managers (who have employees reporting to them) and non-managers:

SQL Query:

```
SELECT DISTINCT Manager_Id AS Manager FROM Employee WHERE Manager_Id IS NOT NULL;
```

```
SELECT Id FROM Employee WHERE Id NOT IN (SELECT DISTINCT Manager_Id FROM Employee WHERE Manager_Id IS NOT NULL);
```

3. Add a new column with the average salary by department

SQL Query:

```
SELECT Id, Name, DEPT, SALARY,
       AVG(SALARY) OVER(PARTITION BY DEPT) AS Avg_Salary_By_Dept
FROM Employee;
```

4. Check if a string is a palindrome

Python Code:

```
def is_palindrome(s):
    return s == s[::-1]

string = "level"
print(is_palindrome(string)) # Output: True
```

5. Duplicate characters in a string ("123a!" → "112233aa!!")

Python Code:

```
def duplicate_chars(s):
    return ''.join([char * 2 for char in s])

string = "123a!"
print(duplicate_chars(string)) # Output: "112233aa!!"
```

6. Check if a number is prime

Python Code:

```
def is_prime(n):
    if n <= 1:
        return False
    for i in range(2, int(n**0.5) + 1):
        if n % i == 0:
            return False
    return True

number = 29
print(is_prime(number)) # Output: True
```



Technical Round 2

1. Processing 1 TB of data in Spark

Steps to efficiently process 1 TB of data:

- Use partitioning and bucketing to distribute data evenly.
- Leverage columnar file formats like Parquet or ORC.
- Optimize shuffle operations with coalesce and reduceByKey instead of groupByKey.
- Utilize broadcast joins for small lookup tables.
- Enable Adaptive Query Execution (AQE) to dynamically optimize the query plan.
- Use caching only for reused datasets.
- Tune executor memory and number of cores based on the cluster size.

2. Designing a Data Warehouse (DWH) Problem Statement

Consider designing a DWH for an e-commerce platform:

- **Fact Tables:**
 - Sales_Fact: Stores transaction details (product_id, customer_id, quantity, revenue, date).
 - Inventory_Fact: Tracks inventory levels by warehouse.
- **Dimension Tables:**
 - Product_Dim: product_id, product_name, category.
 - Customer_Dim: customer_id, name, address.
 - Date_Dim: date_key, day, month, year, quarter.
- **Schema Design:**
 - Use a star schema for easy navigation.
 - Apply partitioning on large tables based on date for efficient querying.
 - Ensure surrogate keys for dimensions.

3. How would you design a data pipeline to handle semi-structured and unstructured data?

- Explain the design using distributed file systems (HDFS, S3), data processing frameworks (Spark, Flink), and NoSQL databases (HBase, MongoDB).
- Include stages: Ingestion (Kafka, Kinesis), Data Lake (Parquet/ORC storage), Transformation (Spark DataFrame, PySpark), and Serving Layer (Redshift, Snowflake, Elasticsearch).
- Discuss schema inference and schema-on-read strategies.

4. Explain the differences between Spark's shuffle and broadcast join. When would you use each?

- **Shuffle Join:**
 - Used when both datasets are large.
 - Requires shuffling of data across nodes.
 - Slower and resource-intensive.
- **Broadcast Join:**
 - Broadcasts a smaller dataset to each executor.
 - Used when one dataset is small enough to fit in memory.
 - Reduces shuffle and improves performance.

When to use:

- Prefer broadcast joins when a small lookup table is involved.
- Use shuffle join for larger datasets, optimizing with partitioning strategies.

5. What strategies can you use to handle skewed data in Spark?

- Salting: Add a random key to distribute data evenly across partitions.
- Increase parallelism: Adjust the number of partitions using repartition().
- Broadcast smaller tables: Use broadcast joins.
- Adaptive Query Execution (AQE): Enable Spark's dynamic partitioning feature.

6. Explain Adaptive Query Execution (AQE) in Spark 3.x.

- AQE optimizes queries dynamically at runtime based on the actual data processed.
- **Key Features:**
 - Dynamic partition pruning: Reduces partitions scanned based on filter conditions.
 - Join optimization: Automatically switches join strategies (broadcast, shuffle) during execution.
 - Coalescing shuffle partitions: Reduces the number of shuffle partitions.

7. How would you optimize a Spark job that takes too long to run in production?

- **Profiling tools:** Use Spark UI to identify bottlenecks.
- **Techniques:**
 - Cache frequently used DataFrames.
 - Use columnar file formats like Parquet/ORC with predicate pushdown.
 - Avoid using collect() or count() on large datasets.
 - Reduce shuffle operations with efficient transformations (reduceByKey over groupByKey).

8. Explain how you would implement a Slowly Changing Dimension (SCD) Type 2 in Spark.

- Use merge operations with Delta Lake or Apache Hudi.
- Steps:
 - Load current data as the target DataFrame.
 - Load incoming data as the source DataFrame.
 - Identify new records, unchanged records, and updated records.
 - Assign start_date, end_date, and current flag to version data.

9. How do you monitor and debug Spark applications in production?

- **Spark UI:** Provides DAG visualization and job stages.
- **Ganglia/Prometheus:** Used for resource monitoring.
- **Event logs:** Enable detailed job logging.
- **Metrics:** Monitor executor memory usage, garbage collection, and shuffle read/write.

10. How do you design a scalable and fault-tolerant data warehouse on a cloud platform?

- **Storage:** Use S3 or Azure Blob Storage for data lakes.
- **Compute:** Spark on EMR or Databricks for processing.
- **Metadata Management:** Hive Metastore or AWS Glue.
- **Query Engine:** Presto, Redshift, or Synapse Analytics.
- **Fault Tolerance:** Leverage **checkpointing** and **retry policies**.



Glassdoor HashedIn by Deloitte Review –

<https://www.glassdoor.co.in/Reviews/HashedIn-by-Deloitte-Reviews-E689428.htm>

HashedIn by Deloitte Careers –

<https://hashedin.com/careers/>

Subscribe to my YouTube Channel for Free Data Engineering Content –

<https://www.youtube.com/@shubhamwadekar27>

Connect with me here –

<https://bento.me/shubhamwadekar>

Checkout more Interview Preparation Material on –

https://topmate.io/shubham_wadekar

