

# LTI Mindtree Data Engineer Interview Guide – Experienced 3+

## Round 1 - Technical

### 1. Project Explanation and Technologies Used

- Describe the projects emphasizing Spark, Hadoop, or Azure for large-scale data processing.
- Discuss tools for ETL, orchestration (Airflow, ADF), and real-time pipelines (Kafka, Spark Streaming).

### 2. Performance Tuning Techniques

- Optimize shuffles: Use appropriate partitioning.
- Use persist()/cache() for reused datasets.
- Leverage broadcast joins for smaller datasets.
- Optimize query plans using explain().

### 3. Accumulator and Broadcast Variables

- Accumulator: Write-only, for aggregating data across tasks (e.g., counters).
- Broadcast Variables: Read-only, for sharing smaller datasets among executors.

### 4. Difference Between SparkSession and SparkContext

- SparkSession: Unified entry point for DataFrame, SQL, and streaming (introduced in Spark 2.0).
- SparkContext: Core entry point for RDD-based operations.

### 5. Difference Between Dataset and DataFrame

- Dataset: Provides compile-time type safety (Scala/Java).
- DataFrame: Schema-based API; simpler for SQL-like operations.

### 6. Spark Session Command

```
from pyspark.sql import SparkSession  
spark = SparkSession.builder.appName("App").getOrCreate()
```

## 7. Command to Read JSON Data and Options

```
df = spark.read.option("multiline", "true").json("path_to_json")
```

## 8. CSV Without Column Names/Schema

```
df = spark.read.option("header", "false").csv("path_to_csv")
df = df.toDF("col1", "col2", "col3") # Assign custom column names
```

## 9. Find 3rd Highest Salary

```
from pyspark.sql.functions import col, dense_rank
from pyspark.sql.window import Window

window_spec = Window.orderBy(col("salary").desc())
df = df.withColumn("rank", dense_rank().over(window_spec))
third_highest = df.filter(col("rank") == 3)
```

## 10. Filter Rows Where Employee Salary > Manager Salary

```
emp_df = df.alias("emp")
mgr_df = df.alias("mgr")

result = emp_df.join(mgr_df, emp_df["mngName"] == mgr_df["Empname"]) \
    .filter(emp_df["salary"] > mgr_df["salary"]) \
    .
```

## 11. Palindrome

```
def is_palindrome(s):
    return s == s[::-1]
print(is_palindrome("madam")) # Example
```

## 12. Spark Submit

```
spark-submit --class <MainClass> --master <ClusterURL> <AppJar>
```

### **13. Memory Tuning**

- Adjust **executor memory** (`--executor-memory`).
- Use **storage levels** (`MEMORY_AND_DISK`).
- Tune **GC settings** for large heaps.

### **14. Created JARs**

Describe building custom JARs for Spark jobs using Maven/SBT.

### **15. Worked with UDFs**

Share examples of custom UDFs using Python or Scala.

### **16. Dynamic Resource Allocation**

Automatically adjusts resources based on workload. Enable with:

`--conf spark.dynamicAllocation.enabled=true`

### **17. Daily Data Volume**

Quantify daily data (e.g., 1TB/day) and its source (e.g., logs, transactions).

### **18. Production Experience**

Discuss deploying Spark jobs and monitoring them using tools like Airflow or YARN.

## Round 2 – Technical

### 1. Difference Between DataFrame and Dataset

**DataFrame**: Untyped, optimized for SQL operations.

**Dataset**: Typed, provides compile-time safety.

### 2. Load CSV from HDFS

```
df = spark.read.option("header", "true").csv("hdfs://path_to_csv")
```

### 3. Syntax for CSV Loading

```
df = spark.read.option("header", "true").csv("hdfs://path_to_csv")
```

### 4. What is Multiline?

**Multiline** option handles JSON files with nested records across multiple lines.

### 5. No Column Names in CSV

```
df = spark.read.option("header", "false").csv("path_to_csv")
df = df.toDF("col1", "col2", "col3") # Assign column names
```

### 6. Case Class and StructType Syntax

```
# Scala Case Class
case class Employee(name: String, salary: Double, dept: String)

# Python StructType
from pyspark.sql.types import StructType, StructField, StringType, DoubleType
schema = StructType([
    StructField("name", StringType(), True),
    StructField("salary", DoubleType(), True),
    StructField("dept", StringType(), True)
])
```

### 7. Partitioning vs. Bucketing

**Partitioning**: Divides data into directories based on keys.

**Bucketing**: Splits data into fixed-size buckets based on hash functions.

## 8. Closure Function

A function that references variables from its enclosing scope.

## 9. Count of Alphabets in String

```
from collections import Counter
s = "I am a data engineer"
counts = Counter(c for c in s.lower() if c.isalpha())
print(counts)
```

## 10. Difference Between List and Tuple

**List:** Mutable, slower.

**Tuple:** Immutable, faster.

## 11. List Comprehension

```
squares = [x**2 for x in range(10)]
```

**Glassdoor LTI Mindtree Review –**

[https://www.glassdoor.co.in/Reviews/LTIMindtree-work-life-balance-Reviews-EI\\_IE8441464.0,11\\_KH12,29.htm](https://www.glassdoor.co.in/Reviews/LTIMindtree-work-life-balance-Reviews-EI_IE8441464.0,11_KH12,29.htm)

**LTI Mindtree Careers –**

<https://www.ltimindtree.com/careers/>

**Subscribe to my YouTube Channel for Free Data Engineering Content –**

<https://www.youtube.com/@shubhamwadekar27>

**Connect with me here –**

<https://bento.me/shubhamwadekar>

**Checkout more Interview Preparation Material on –**

[https://topmate.io/shubham\\_wadekar](https://topmate.io/shubham_wadekar)